



HPC Job Submission

HPC Job Submission

Daniel Trahan

Daniel.Trahan@Colorado.EDU

rc-help@colorado.edu

Slides available for download at

https://github.com/ResearchComputing/Supercomputing_Spin_Up_Fall_2021

Adapted from presentations by RC members Andrew Monaghan, Aaron Holt and John Blaas: [1](#), [2](#), [3](#), [4](#).

Outline

- General Info
- Examples of submitting jobs to the supercomputer!
 - Simple batch jobs
 - Advanced batch jobs: running programs, mpi
 - Interactive jobs

RMACC Summit Supercomputer

- 450 compute nodes (Intel Xeon Haswell)
 - 24 cores per node
 - 11,400 total cores
 - Omni-Path network
 - 1.2 PB scratch storage
 - GPFS File system
-
- 67% CU, 23% CSU, 10% RMACC



Additional Types of RMACC Summit Compute Nodes

- 10 Graphics Processing Unit (GPU) Nodes
 - NVIDIA Tesla K80 (2/node)
- 5 High Memory Nodes
 - 2 TB of memory/node, 48 cores/node
- 20 Phi Nodes
 - Intel Xeon Phi
 - 68 cores/node, 4x threads/core

RC Access: Logging in

- If you have an RMACC account already, login as follows from a terminal:

```
$ ssh <username>@login.rc.colorado.edu  
# Where username is your identikey
```

- If you do not have an RMACC account use one of our temporary accounts:

```
$ ssh user<XXXX>@tlogin1.rc.colorado.edu  
# Where user<XXXX> is your temporary username
```

Working on RC Resources

- When you first log in, you will be on a login node. Your prompt will look like this (e.g.):

```
[user0049@tlogin1 ~]$
```

- The login nodes are lightweight virtual machines primarily intended to serve as 'gateways' to RC resources. In order to get a better view of the software available on Summit we will go to a compile node.

```
[user0049@tlogin1 ~]$ ssh scompile
```

- Now go to your working directory and download the material for this workshop:

```
[user0049@shas0137 ~]$ cd /scratch/summit/$USER  
[user0049@shas0137 ~]$ git clone  
https://github.com/ResearchComputing/Supercomputing_Spin_Up_Fall_2021
```

Jobs

- Because Summit is a shared resource with many users trying to utilize available compute with their applications, we need a system to divide compute in a simple and fair system.
- SLURM
 - **S**imple **L**inux **U**tility for **R**esource **M**anagement
 - (Actually named after the Futurama Drink...)
- Through SLURM, users can grab allotments of compute resources called Jobs
- 2 Types of Jobs
 - Batch Jobs
 - Interactive Jobs

Batch Jobs

- Batch Jobs are jobs you submit to the scheduler where they are run later without supervision.
 - By far the most common job on Summit
 - "batch of cookies"
 - Requires a job script
- A job script is simply a script that includes SLURM directives ahead of any commands.

Submitting a batch job

- `sbatch`: submit a batch job
- Submit your first job! :

```
$ sbatch submit_test.sh
```

- Script contains most of the parameters needed to define a job
- Additional flags can be used to temporarily replace any set parameters.

```
$ sbatch --reservation=supercomputing-spinup-pt2 submit_test.sh
```

<http://slurm.schedmd.com/sbatch.html>

Anatomy of a job script (submit_test.sh)

```
#!/bin/bash
#SBATCH --ntasks=1           # Number of requested tasks
#SBATCH --time=0:01:00       # Max wall time
#SBATCH --partition=shas-testing # Specify Summit Haswell nodes
#SBATCH --output=test_%j.out  # Rename standard output file

# Updated by:  Shelley Knuth, 17 May 2019
# Purpose:    To demonstrate how to run a batch job on RC resources

# Purge all existing modules
module purge

# Run commands
echo "This is a test of user $USER"
```

Job Options

Specified at command line or in job script as...

#SBATCH <options> ...where options include:

- **Partition:** `--partition=<partition_name>`
- Sending emails: `--mail-type=<type>`
- Email address: `--mail-user=<user>`
- Number of nodes: `--nodes=<nodes>`
- Number of cores: `--ntasks=<number-of-tasks>`
- **Quality of service:** `--qos=<qos>`
- **Allocation:** `--account=<account_name>`
- Wall time: `--time=<wall time>`
- Job Name: `--job-name=<jobname>`
- Output: `--output=<name>`

More on slurm commands:

<https://slurm.schedmd.com/quickstart.html>

FYI: You do NOT actually type <> above – this designates something specific you as a user must enter about your job

Partitions

- Partitions specify the type of compute node that you wish to use
 - Specify with the `--partition` flag

```
#SBATCH --partition=shas
```

Partition	Description	# of nodes	cores/node	GPUs/node
shas	General Compute (Haswell)	~450	24	0
sgpu	GPU-enabled nodes	10	24	effectively 4
smem	High-memory nodes	5	48	0
sknl	Phi (Knights Landing) nodes	20	68	0

Sub-Partitions

- In addition to normal compute partitions, Summit Users also have access to several testing and interactive partitions
 - Quick access to get your applications functional!

Partition	Description	Max wall time	Max jobs/user	Max nodes/user
shas-testing sgpu-testing sknl-testing	For quick turnaround when testing	30 M	1	2 12 cores/node
shas-interactive	For interactive jobs (command or GUI)	4 H	1	1 core

Quality of Service (--qos)

- Quality of Service specifies additional constraints Job
 - On Summit, this means if your job needs to run longer than 1 day
 - Specify with the `--qos` flag
 - Doesn't need to be set otherwise

```
#SBATCH --qos=long
```

QoS	Description	Max wall time	Max jobs/user	Max nodes/user
normal	Default QoS	Derived from partition	n/a	256
long	For jobs needing longer wall times	7 D	n/a	20

Example 1: Writing your first job script

Write your first job script!

- Create a job script and submit it as a batch job with the following instructions:
 1. Navigate to the 'job_submission' directory
 2. Name it 'submit_sleep.sh'
 3. The job should contain the following commands:

```
echo "Running on host" `hostname`  
echo "Starting Sleep"  
sleep 30  
echo "Ending Sleep. Exiting Job!"
```

Details on job script parameters are in the next slide

Job details of `submit_sleep.sh`

1. The job will run on 1 core of 1 node
2. We will request a 1 minute wall time
3. Run on the shas-testing partition
4. Set the output file to be named “sleep.%j.out”
5. Name your job “sleep”
6. Bonus: Email yourself when the job ends
7. Contains the following commands →

```
echo "Running on host" `hostname`  
echo "Starting Sleep"  
sleep 30  
echo "Ending Sleep. Exiting Job!"
```

```
$ sbatch --reservation=supercomputing-spinup-pt2 submit_sleep.sh
```

Solution are prefixed with 'answer'

Job Output

- Once a job completes its execution, then standard output of the script will be redirected to an output file.
 - Great for debugging!
 - Could be different from output generated by your application
 - File is created in directory job was run unless specified in your `--output` directive.
 - If the directive `--output` is not provided then a generic file name will be used.

```
$ cat sleep.xxxxxx.out # where xxxxxx is your Job Id
```

Solution can be found in “./solutions” subdirectory

Checking your jobs (1)

- **squeue**: Monitor your jobs status in queue and while running:
 - By Default shows all jobs in queue
 - Narrow this down with:

```
$ squeue -u <username>  
$ squeue -p <partition>
```

- **sacct**: Check back on usage statistics of previous Jobs
 - By default only checks all jobs from the start of the current day.
 - Narrow this down with:

```
$ sacct -u <username>  
$ sacct --start=MM/DD/YY -u <username>  
$ sacct -j <job-id>
```

More on slurm commands: <https://slurm.schedmd.com/quickstart.html>

Checking your jobs (2)

- Another method of checking details of your job is with `scontrol`
- Advanced command usually used by system administrators, but you can use it too!

```
$ scontrol show job <job number>
```

More on slurm commands: <https://slurm.schedmd.com/quickstart.html>

Software and Jobs

- "Okay so running a job is easy, but how do I run a job with my software?"
- LMOD
 - Module system on CURC systems
 - Modifies your environment to make your desired software visible to your terminal.

```
$ module load matlab  
$ ml matlab #shorthand version!
```

More on slurm commands: <https://slurm.schedmd.com/quickstart.html>

Software and Jobs (2)

- More LMOD commands:

```
$ module purge           #Unloads all current modules
$ module unload matlab   #Unloads matlab
$ module spider matlab    #Searches for matlab in module tree
```

- What if my software isn't available through LMOD?
 - Software must be installed locally if not available through LMOD
 - RC User support is happy to assist, but *installs are best effort*
 - For more assistance contact rc-help@colorado.edu

More on slurm commands: <https://slurm.schedmd.com/quickstart.html>

Running an external program

- Let's run R on an R script
- This script calls and runs the script *R_program.R*
- Let's examine the batch script [submit_R.sh](#)
 - Note how R is loaded
- Go ahead and submit the batch script [submit_R.sh](#)

Example 2: Serial Matlab Code

Launch Matlab!

- Create a job script and submit it as a batch job with the following instructions:
 1. Name it 'submit_matlab.sh'
 2. Load the 'matlab' module
 3. The job should contain the following commands:

```
cd progs  
matlab -nodisplay -nodesktop -r "matlab_tic;"
```

Details on job script parameters are in the next slide

Job details of `submit_matlab.sh`

1. The job will run on 1 core of 1 node
2. We will request a 2 minute wall time
3. Run on the shas-testing partition
4. Set the output file to be named “matlab.%j.out”
5. Name your job “matlab”
6. Bonus: Email yourself when the job ends
7. Contains the following commands →

```
cd progs  
matlab -nodisplay -nodesktop -r "matlab_tic;"
```

```
$ sbatch --reservation=supercomputing-spinup-pt2 submit_sleep.sh
```

Solution are prefixed with 'answer'

Running an mpi job

- For cases where you have a code that is parallelized, meaning it can run across multiple cores.
- Number of tasks always > 1 . E.g.,

```
#SBATCH --ntasks=4
```

- Will always need to load a compiler and mpi. E.g.,

```
module load intel impi
```

- Executable preceded with mpirun, srun, or mpiexec. E.g.,

```
mpirun -np 4 python yourscrip.py
```

- Examine and run the example '[submit_python_mpi.sh](#)'

```
$ sbatch --reservation=supercomputing-spinup-pt2 submit_python_mpi.sh
```

Interactive jobs

- Sometimes we want our job to run in the background
- Sometimes we want to work on program in real time
 - Great for testing, debugging
- For example, let's run the R job we previously ran as a batch job, but this time let's do it interactively...

Running an interactive job

- To work with R interactively, we request time from Summit
- When the resources become available the job starts
- Commands to run:

```
$ sinteractive --time=00:10:00 --reservation=hpc
```

- Once we receive a prompt, then:

```
$ module load R  
$ cd ./progs  
$ Rscript R_program.R
```

- Once we finish we must exit! (job will time out eventually)

```
$ exit
```

Tools for submitting many small jobs at once

- CURC Load Balancer
 - <https://curc.readthedocs.io/en/latest/software/loadbalancer.html>
- Slurm job arrays
 - https://slurm.schedmd.com/job_array.html

Thank you!

- Please fill out the survey: <http://tinyurl.com/curc-survey18>
- Contact information: rc-help@Colorado.edu
- Slides and Examples from this course:
https://github.com/ResearchComputing/Supercomputing_Spin_Up_Fall_2021
- Slurm Commands: <https://slurm.schedmd.com/quickstart.html>